

Verifying computations without re-executing them

The GKR Protocol

Myungsun Kim

Department of Information Security
The University of Suwon

September 27, 2019

@Hanyang University

Contents

1 Introduction

Motivations

High-level overview

2 The GKR Protocol

Warming-up

The Sum-Check Protocol

The GKR Protocol

1. Introduction

§1.1 Motivations

A classic problem in computation theory

The question

How can a single PC check a herd of supercomputers with unreliable software and untested hardware?

[BFLS91]

In this setup, a single reliable PC can monitor the operation of a herd of supercomputers working with possibly extremely powerful but unreliable software and untested hardware.

– Babai, Fortnow, Levin, and Szegedy: *Checking computations in polylogarithmic time*. STOC, 1991.

Due to Babai et al.

In an asymmetric setting (a PC vs. dozens of Supercomputers)
with available resources,

Delegating some expensive tasks to the cloud:

- Pros
 - ▶ Fixed cost & depreciation savings
 - ▶ Supplies expenses savings
 - ...
- Cons
 - ▶ Black-box in faults: mis-config., corruption of data in transit, H/W problems, & mal. op.'s
 - ...

A central issue

Is an output from the cloud as a response to a service request correct?

Due to Babai et al.

Possible answers

- ① Replicate of delegations
- ② Audit: checking the responses in a small sample
- ③ Trust H/W (a.k.a TPM) & attestation

...

Another possible solution—(★)

- Supercomputers return some results along with a proof that the results were correctly computed
- Checking the proof is cheap (vs. locally redoing the computation)

A bit more realistic scenario

Consider a simple DB with a table [EN11], that

- keeps track of employee's name, social security number, address, salary, and her working department

EMPLOYEE


Fname	Lname	Ssn	Address	Salary	Dno
John	Smith	1234	73 Fondren, Houston, TX	30000	5
Franklin	Wong	3334	68, Voss, Houston, TX	40000	5
Alicia	Zelaya	9998	21, Castle, Spring, TX	25000	4
Jennifer	Wallace	9876	91, Berry, Bellaire, TX	43000	4
Ramesh	Narayan	6688	75, Oak, Humble, TX	38000	5
Joyce	English	4534	56, Rice, Houston, TX	25000	5
Almad	Jabbar	9879	80, Dallas, Houston, TX	31000	4
James	Borg	8886	50, Stone, Houston, TX	55000	1

A bit more realistic scenario

A standard DBMS provides DB query services

- Search (e.g., `select-from-where`)
- Update (e.g., `insert or delete`)

by running the DBMS



A bit more realistic scenario

When the client submits a query **x** to the DBMS

```
1  SELECT Fname, Lname, Salary, Address  
2  FROM EMPLOYEE  
3  WHERE Salary = (SELECT MAX(E1.Salary)  
4                      FROM EMPLOYEE E1, PROJECT P  
5                      WHERE P.Pname = "Design"  
6                      AND P.Dnum = E1.Dno  
7                      AND E1.Age >= 37)  
8  ORDER BY Fname, Lname;
```


Z = +-----+
| Fname | Lname | Salary | Address |
+-----+
| John | Smith | 30000 | 73, Fondren, Houston, TX |
| Joyce | English | 25000 | 56, Rice, Houston, TX |
+-----+

A bit more realistic scenario

Main security concerns

- ① **Correctness:** $z = f_y(\mathbf{x})?$
 - ② **Privacy:** The DBMS learns my secret data \mathbf{x}, y and/or z
-
- 😊 From homomorphic encryption (E, D) (e.g., [CKK17]):
 - ▶ encrypt a DB y into $\bar{y} \leftarrow E(y)$
 - ▶ store them to the remote server;
 - ▶ encrypt a query \mathbf{x} into $\bar{\mathbf{x}} \leftarrow E(\mathbf{x})$ and send $\bar{\mathbf{x}}$ to the server
 - ▶ evaluate f at $\bar{\mathbf{x}}, \bar{y}$ as \bar{z}
 - ⌚ How to verify if $z^* \leftarrow D(\bar{z})$ is a correct query result?
 - ▶ $z^* = f_y(\mathbf{x})?$
 - ▶ E.g., vSQL in [ZGK+17], but over clear DB

A short history in verifiable computation




§1.2 The Abstract Model

A system model

A verifiable computation (VC) system:


- Goal: Proving integrity to the verifier,
 - ▶ Prover's auxiliary input is secret \Rightarrow privacy-preserving integrity
- Participating players
 - ▶ Verifier: Wish to know the output from a computation p on an input x
 - ▶ Prover: Wish to efficiently convince the verifier of his computing the output y , while returning y and a certificate of correctness π
- Adversarial model
 - ▶ The verifier is assumed to be honest, but the prover is dishonest

A way to enforce security



A way to enforce security

Applying to our scenario



A way to enforce security

The key to the powerfulness of  (the mechanic before) is the **combination** of V 's **randomness** and (P, V) 's **interaction**, for the purpose of amplifying soundness guarantee by

- Tossing random coins: even a supercomputer does not know but just guess
 - ▶ The distinction bet. public & internal coins is not something [GS86]
- Interacting many times: reduce the small probability to be negligible

The system model

Interactive proof (IP) systems : Informal

- Participants
 - ▶ Prover P (modeling the DBMS or the suspect)
 - ▶ Verifier V (modeling the client or the interrogator)
- Activities
 - ① P solves a problem on a given input, **modeling the question**
“Is the result z correct?”
 - ② P shows the answer
 - ③ P **proves** to V that the answer is correct
- Security requirements
 - ▶ **Completeness** states that an honest P always can convince V to accept
 - ▶ **Soundness** states that V is able to catch a cheating P with high probability

Interactive Proof Systems

Classifications of proof systems:

- Interactive proofs (IPs) [GMR89,Bab85]: ensure information-theoretically soundness
- Argument systems [BCC88]: ensure soundness just against polynomial time P 's
- Multi-prover IPs (MIPs)
 - ▶ MIP [BGKW88]: similar to IP, but there are 2 or more P 's
Note. In PCP, the answer (formally a proof) is static, but may have super-polynomial size
- Zero-knowledge proofs and its argument variants: ensure that no information are revealed to V except the validity of answer being given

Interactive Proof Systems

Zero-knowledge proofs (ZKPs)

- Some IP and argument systems \Leftrightarrow ZKPs (or ZK arguments)
 - ▶ P with an ingenious know-how to persuade V
 - ▶ P in such IPs (and arguments) reveals no knowledge to V other than the correctness of his proof
- Additional applications from IPs of achieving ZK
 - ▶ The **sum-check** protocol [LFKN92]: a building block in efficient delegating computation
 - during execution, reveals to V partial sums over V 's randomness
 - the partial sums $\not\Rightarrow$ ZK-**sum-check**
 - ⌚ with higher costs than expected

Definition

A formal definition of IP

Definition. For a function $f : \{0, 1\}^n \rightarrow R$, an interactive proof system for f consists of a PPT verifier V and a prover P with a common input $\mathbf{x} \in \{0, 1\}^n$. After poly. many comm.'s, denoted by $\mathbf{t} := (V(\mathbf{r}), P)(\mathbf{x})$, V outputs $\{0, 1\}$. The IP system has δ_c, δ_s if the two conditions hold.

- ① (Completeness) $\exists P$ such that $\forall \mathbf{x} \in \mathcal{L}$:

$$\Pr[\text{out}_V(V(\mathbf{r}), P)(\mathbf{x}) \rightarrow 1] \geq 1 - \delta_c$$

- ② (Soundness) $\forall \mathbf{x} \neq \mathcal{L}$ and P^* :

$$\Pr[\text{out}_V(V(\mathbf{r}), P^*)(\mathbf{x}) \rightarrow 1] \leq \delta_s$$

where \mathbf{r} is V 's internal random; the IP for f is equivalent to the IP for language $\mathcal{L}_f = \{(\mathbf{x}, y) | y \leftarrow f(\mathbf{x})\}$

2. The GKR Protocol

§2.1 Warming-up

A warmup example

Freivalds' algorithm

- For $A, B \in \mathcal{M}_{n \times n}$ over \mathbb{F}_p , the fastest algorithm to compute $A \cdot B$ runs in $O(n^{2.3728639})$ [LG14]
- The Freivalds algorithm [Fre77]
 - 1 P opens $C \leftarrow A \cdot B$
 - 2 V chooses $r = (r, r^2, \dots, r^n)$ where $r \in \mathbb{F}_p$
 - 3 Compute $y \leftarrow Cr$ and $z \leftarrow A \cdot (Br)$
 - 4 Check $y = z$; if yes, accept, reject otherwise
- V of the Freivalds algorithm runs in $O(n^2)$
- If $A \cdot B \neq C$, then for some i : $C_i \neq (A \cdot B)_i$ and

$$\Pr[(Cr)_i = (A \cdot (Br))_i] \leq n/p$$

A warmup example

Analysis

Fact

For any two distinct polynomial p_a and p_b of degree at most n with coefficients in \mathbb{F}_p , $p_a(x) = p_b(x)$ for at most n values of x in \mathbb{F}_p where $p_a(x) = \sum a_i x^i$.

- Let $D = A \cdot B$ and observe only the case $\exists i$ s.t. $C_i \neq D_i$
- View $(Cr)_i = p_{C_i}(r)$ and $A \cdot (B \cdot r) = p_{D_i}(r)$
- $\Pr[(Cx)_i \neq A \cdot (B \cdot x)_i] \geq 1 - \frac{n}{p}$
- Apply the fact

§2.3 The Sum-check Protocol

The sum-check protocol

- Goal: Given a v -variate polynomial g over \mathbb{F} , to compute the sum

$$H := \sum_{x_1 \in \{0,1\}} \sum_{x_2 \in \{0,1\}} \cdots \sum_{x_v \in \{0,1\}} g(x_1, x_2, \dots, x_v)$$

- The strategy [LFKN92]
 - ▶ $H = \sum_{x_2 \in \{0,1\}} \cdots \sum_{x_v \in \{0,1\}} g(0, x_2, \dots, x_v) + \sum_{x_2 \in \{0,1\}} \cdots \sum_{x_v \in \{0,1\}} g(1, x_2, \dots, x_v)$
 - ▶ Each partial sum is the evaluation of $g(t, x_2, \dots, x_v)$ at 0 or 1
 - ▶ Proceed evaluations incrementally while binding the variable t_j to a random $r_j \in \mathbb{F}$

The sum-check protocol

- ① Fix an $H \in \mathbb{F}$
- ② 1st-round.
 - ▶ P sends the univariate polynomial

$$g_1(t_1) = \sum_{x_2, \dots, x_v \in \{0,1\}^{v-1}} g(t_1, x_2, \dots, x_v)$$

- ▶ V checks if $g_1(t_1)$ has the claimed degree and $H = g_1(0) + g_1(1)$; if not rejecting
- ▶ V sends a random $r_1 \in \mathbb{F}$ to P

The sum-check protocol

③ j^{th} -round ($1 < j < v$)

- ▶ P sends the univariate polynomial

$$g_j(t_j) = \sum_{x_{j+1}, \dots, x_v \in \{0,1\}^{v-j}} g(r_1, r_2, \dots, r_{j-1}, t_j, x_{j+1}, \dots, x_v)$$

- ▶ V checks if $g_j(t_j)$ has the claimed degree and $g_{j-1}(r_{j-1}) = g_j(0) + g_j(1)$; if not rejecting
- ▶ V sends $r_j \xleftarrow{\$} \mathbb{F}$ to P

The sum-check protocol

The sum-check protocol

4 v^{th} -round.

- ▶ P sends the univariate polynomial

$$g_v(t_v) = g(r_1, \dots, r_{v-1}, t_v)$$

- ▶ V checks if $g_v(t_v)$ has the claimed degree and $g_{v-1}(r_{v-1}) = g_v(0) + g_v(1)$; if not rejecting
- ▶ V picks $r_v \xleftarrow{\$} \mathbb{F}$ and check if $g_v(r_v) = \underbrace{g(r_1, r_2, \dots, r_v)}_{\text{How efficiently?}}$; if not rejecting

5 V halts or accepts if it has not yet rejected

The sum-check protocol

Efficiency of the SC protocol

- Costs of the sum-check protocol

Communication	Rounds	V's time	P's time
$O(\sum_{i=1}^v \deg(g_i)) \mathbb{F} \text{ elt's}$	v	$O(\sum_{i=1}^v \deg(g_i)) + T$	$O(2^v \cdot T)$

- Soundness error $\leq \frac{dv}{|\mathbb{F}|}$ by induction on v & the Fact

§2.4 The Main Protocol

An example


Consider a piece of a C code where x_1, x_2, y are bits

```
/* y= x1 ^ x2 */
void foo(const char* x)
{
    if (x1 != x2) {
        y = 1;
    }
    else {
        y = 0;
    }
    /* do something more */
    return y;
}
```

An example

Representing the program `foo` as a circuit in terms of AND, OR and NOT

- $y = x_1 \oplus x_2$
- $y = (\neg x_1 \wedge x_2) \vee (x_1 \wedge \neg x_2)$
- A circuit \mathcal{C}_{foo}



Background

Some background: Low degree & multilinear extensions

Notation:

\mathbb{F} is a finite field, and $f : \{0, 1\}^v \rightarrow \mathbb{F}$ is a function

Definition. A polynomial $g \in \mathbb{F}[x_1, \dots, x_v]$ is said to be an extension of f if $\forall \mathbf{x} \in \{0, 1\}^v : g(\mathbf{x}) = f(\mathbf{x})$.

Definition. A multivariate polynomial g is multilinear if the degree of g in each variable ≤ 1 .

Lemma. Any function $f : \{0, 1\}^v \rightarrow \{0, 1\}$ has a unique multilinear extension (MLE) over \mathbb{F} , denoted by \tilde{f}

Proof.

Background

Usefulness of low-degree extensions.

- An MLE g of f : a distance-amplifying encoding of f
- $f \neq f'$ at a single input $\Rightarrow \tilde{g} \neq \tilde{g}'$ almost everywhere

Example: $f : \{0, 1\}^2 \rightarrow \mathbb{F}_5$ and $f(0, 0) = 1, f(0, 1) = 2,$
 $f(1, 0) = 1, f(1, 1) = 4$

- $\tilde{f}(x_1, x_2) = 2x_1x_2 + x_2 + 1$

$$\begin{array}{cc} & \begin{matrix} 0 & 1 \end{matrix} \\ \begin{matrix} 0 & \\ 1 & \end{matrix} & \xrightarrow{\text{extension}} \end{array}$$

f	1	2
1	1	4

	0	1	2	3	4
0	1	2	3	4	0
1	1	4	2	0	3
2	1	1	1	1	1
3	1	3	0	2	4
4	1	0	4	3	2

$$g = \tilde{f}$$

The GKR Protocol

Protocol overview:

- Goal: Compute the value of the output gate(s) of \mathcal{C} for a logspace arithmetic circuit \mathcal{C} whose size S and depth d over n variables
- Costs of the GKR protocol [GKR08]

Communication	Rounds	V 's time	P 's time
$d \cdot \text{polylog}(S)$ field elt's	$d \cdot \text{polylog}(S)$	$O(n + d \cdot \text{polylog}(S))$	$\text{poly}(S)$

- Protocol activities:
 - ① P sends to V the claimed output of \mathcal{C}
 - ② Bind a claim about level $i \geq 1$ and a claim about level $i + 1$ by V 's random
 - ③ At level $d + 1$, V locally check if the last claim is consistent with the public input x

The GKR Protocol

A bit more details of the GKR protocol

- At the first round, V is given a claim about the value(s) if the output gate(s) of \mathcal{C}
 - :(V cannot check this claim by herself)
- V reduces the claim about the outputs of \mathcal{C} to a claim about the gate values at level 2,
 - :(using the sum-check protocol)
- Do the same thing by level d
- At the $(d + 1)$ level, V is given a claim about the inputs \mathbf{x} of \mathcal{C}
 - :(V locally checks the claim using the inputs \mathbf{x})

The GKR Protocol

Notation

- \mathcal{C} : a layered arith. circuit of size s , depth d , and fan-in 2 where 1: the output level & $d + 1$: the input level
- s_i : the number of gates at level i and $s_i = 2^{\sigma_i}$ for some $\sigma_i \geq 0$
- $W_i : \{0, 1\}^{s_i} \rightarrow \mathbb{F}$: take a binary gate label, output the corr. gate's value \widetilde{W}_i which is W_i 's MLE version
- $\text{in}_{1,i}, \text{in}_{2,i} : \{0, 1\}^{s_i} \rightarrow \{0, 1\}^{s_{i+1}}$: take the label a of a gate at level i , output the label of each in-neighbor of gate a
- $\text{add}_i : \{0, 1\}^{s_i+2s_{i+1}} \rightarrow \{0, 1\}$: a wiring predicate with its MLE $\widetilde{\text{add}}_i$
- $\text{mul}_i : \{0, 1\}^{s_i+2s_{i+1}} \rightarrow \{0, 1\}$: a wiring predicate with its MLE $\widetilde{\text{mul}}_i$ where a wiring predicate **encodes**
 - ▶ which pairs of wires from level $i + 1$ are connected to
 - ▶ a given gate at level i in \mathcal{C}

The GKR Protocol

Details of the GKR protocol

- with d iterations;
- each iteration i starts with P claiming that a value for $\widetilde{W}_i(\mathbf{r}_i)$ for V 's random $\mathbf{r}_i \in \mathbb{F}^{s_i}$

① 1-round: check the claim about the outputs ($s_0 = 2^{\sigma_0}$) in \mathcal{C}

- ▶ V
 - uses $D : \{0, 1\} \rightarrow \mathbb{F}$, a function mapping the label of the output gate to the claimed value of that output
 - picks $\mathbf{r}_0 \in \mathbb{F}^{\sigma_0}$ and evaluates $\widetilde{D}(\mathbf{r}_0)$ in time $O(\sigma_0)$ [VSBW13]
 - checks if $\widetilde{D}(\mathbf{r}_0) = \widetilde{W}_0(\mathbf{r}_0)$, where
 $\widetilde{D}(\mathbf{r}_0) = \widetilde{W}_0(\mathbf{r}_0)$: the MLE of claimed output = the MLE of the corrected outputs when evaluated at a random point \mathbf{r}_0
- ▶ V can evaluate $\widetilde{W}_0(\mathbf{r}_0)$ only with P 's interactions

The GKR Protocol

② i -round ($i \leq d$):

- ▶ Goal: reduce the claim about the value of $\widetilde{W}_i(\mathbf{r}_i)$ to a claim about $\widetilde{W}_{i+1}(\mathbf{r}_{i+1})$ for V 's random $\mathbf{r}_{i+1} \in \mathbb{F}^{\sigma_{i+1}}$
- ▶ Run the sum-check protocol to a polynomial $f_{\mathbf{r}_i}^{(i)}$ derived from \widetilde{W}_{i+1} , $\widetilde{\text{add}}_i$ and $\widetilde{\text{mul}}_i$
- ▶ To check P 's claim about $\widetilde{W}_i(\mathbf{r}_i)$:
 - (a) Express $\widetilde{W}_i(\mathbf{r}_i)$ as

$$\begin{aligned} W_i(\mathbf{a}) = \sum_{\mathbf{b}, \mathbf{c} \in \{0,1\}^{\sigma_{i+1}}} & \widetilde{\text{add}}_i(\mathbf{a}, \mathbf{b}, \mathbf{c})(\widetilde{W}_{i+1}(\mathbf{b}) + \widetilde{W}_{i+1}(\mathbf{c})) + \\ & \widetilde{\text{mul}}_i(\mathbf{a}, \mathbf{b}, \mathbf{c})(\widetilde{W}_{i+1}(\mathbf{b}) \cdot \widetilde{W}_{i+1}(\mathbf{c})) \end{aligned}$$

- (b) Apply to the sum-check protocol to $f_{\mathbf{r}_i}^{(i)}$

$$\begin{aligned} f_{\mathbf{r}_i}^{(i)}(\mathbf{b}, \mathbf{c}) = & \widetilde{\text{add}}_i(\mathbf{r}_i, \mathbf{b}, \mathbf{c})(\widetilde{W}_{i+1}(\mathbf{b}) + \widetilde{W}_{i+1}(\mathbf{c})) + \\ & \widetilde{\text{mul}}_i(\mathbf{r}_i, \mathbf{b}, \mathbf{c})(\widetilde{W}_{i+1}(\mathbf{b}) \cdot \widetilde{W}_{i+1}(\mathbf{c})) \end{aligned}$$

The GKR Protocol

③ The final round:


► V

- The vector of gate values at level $d + 1$ is the input \mathbf{x} of \mathcal{C}
- Locally evaluate $\mathbf{W}_{d+1}(\mathbf{r}_{d+1})$ in time $O(n)$

The GKR Protocol

Costs of the GKR protocol

- V 's time
 - ▶ Run the sum-check to evaluate $f_{\mathbf{r}_i}^{(i)}(\mathbf{b}, \mathbf{c})$ which is $2\sigma_{i+1}$ -variate & of deg. at most 2 in each var. $\Rightarrow 2\sigma_{i+1}$ -round at level i
 - ▶ the total time cost is $O(n + d \log S + m)$ where
 - processing time $d \log S$ messages between V and P
 - m for evaluating $\widetilde{\text{add}}_i, \widetilde{\text{mul}}_i$ & n for evaluating $\mathbf{W}_{d+1}(\mathbf{r}_{d+1})$
- P 's time: d times running of the sum-check protocol for MLEs

A large, faint watermark of the University of Suwon logo is visible in the background. The logo is circular with the text "UNIVERSITY OF SUWON" around the top and "The University of Suwon" in Korean ("수원대학교") at the bottom.

Thank you & Questions?